BNPL Credit Risk — Modeling (Logistic Regression and k-NN)

ISOM3018 Course Project

Contents

Section 1: Modeling Overview **1**

Section 2: Logistic Regression **2**

Section 3: k-Nearest Neighbors (k-NN) **4**

Section 4: Model Comparison **6**

Section 5: Business Implications **7**

Section 1: Modeling Overview

Two classification models — logistic regression and k-nearest neighbors (k-NN) — are built to predict whether a BNPL customer will default on their repayment (`default_flag`: Default vs. No Default).

Both models are trained on the same preprocessed training set (SMOTE applied to address class imbalance) and evaluated on a held-out test set. Numeric predictors are standardized before fitting. The positive class for all discrimination metrics is **Default**.

Section 2: Logistic Regression

Logistic regression models the probability of default as a linear function of customer characteristics. It is selected as a baseline for its interpretability — the contribution of each predictor can be directly examined. All numeric predictors are standardized; categorical predictors are dummy-encoded.

Table 1: Test Set Performance — Logistic Regression

Metric	Value
roc_auc	0.7439
accuracy	0.6607
sens	0.6473
spec	0.6693

Interpretation: The model achieves a ROC AUC of 0.74, indicating moderate discriminative ability. At the default 0.50 threshold, it correctly identifies around 65% of actual defaulters (sensitivity), while maintaining a similar rate for non-defaulters (specificity). For a BNPL provider, higher sensitivity is generally preferred — missing a defaulter is more costly than a false alarm.

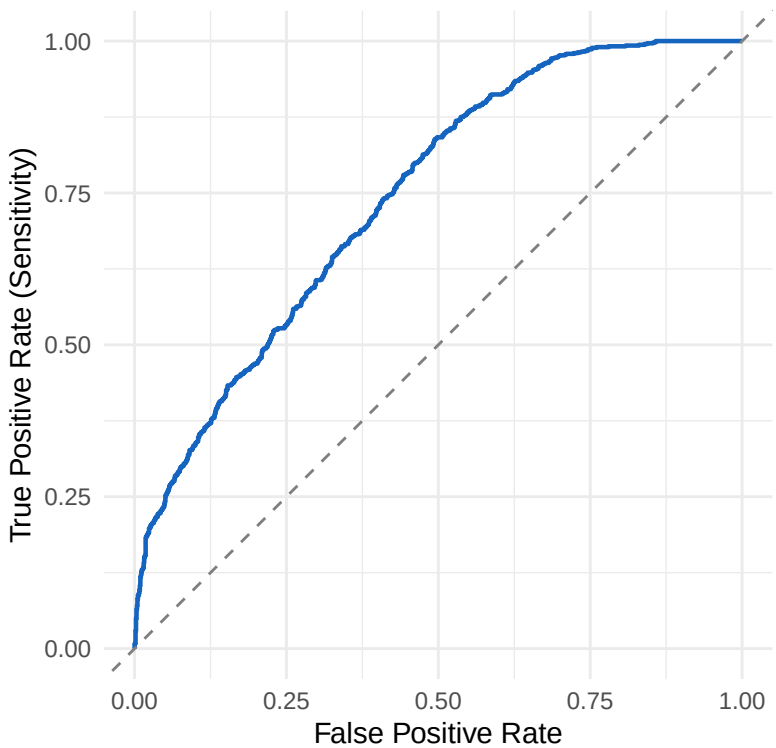

Figure 1: ROC Curve — Logistic Regression (Test Set).

Interpretation: The ROC curve lies clearly above the diagonal (random baseline), confirming the model has meaningful predictive power for identifying BNPL default risk.

Table 2: Confusion Matrix — Logistic Regression (Test Set)

	Predicted: Default	Predicted: No Default

Actual: Default	523	417
Actual: No Default	285	844

Interpretation: Of the 808 customers who actually defaulted, 523 were correctly flagged (true positives), while 285 were missed (false negatives). Of the 1,261 non-defaulters, 844 were correctly cleared and 417 were incorrectly flagged. Reducing the false negative count — undetected defaulters — is the primary operational priority for credit risk management.

Section 3: k-Nearest Neighbors (k-NN)

k-NN classifies each customer based on the majority default outcome among their k nearest neighbors in standardized feature space. The optimal k is selected by 5-fold cross-validation on the training data, choosing the value that maximizes ROC AUC.

Table 3: Optimal k Selected by Cross-validation

Selected k	Selection Criterion
55	Max CV ROC AUC

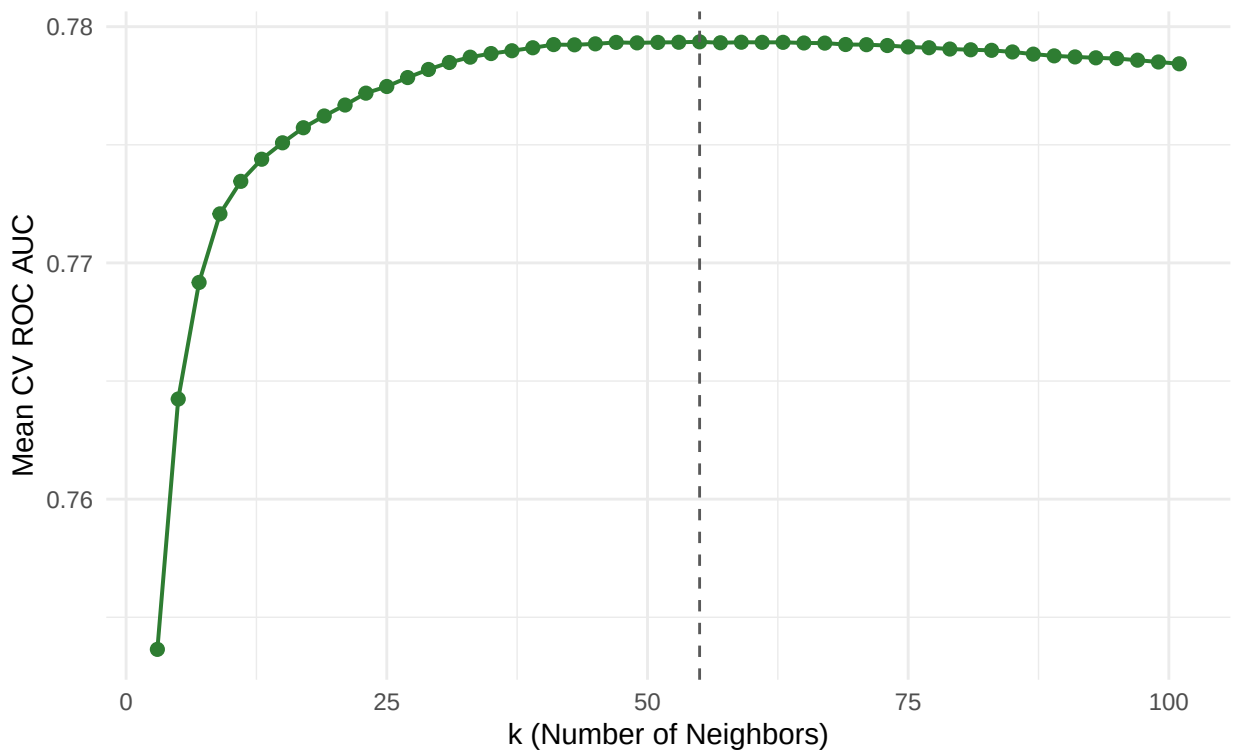

Figure 2: Cross-validated ROC AUC by k (Training Data Only).

Interpretation: ROC AUC increases steadily as k grows, peaking at k = 55. A larger k produces smoother, more stable predictions by averaging over more neighbors, which helps on this dataset.

Table 4: Test Set Performance — k-NN

Metric	Value
roc_auc	0.7289
accuracy	0.6477
sens	0.6609
spec	0.6392

Interpretation: k-NN achieves a ROC AUC of 0.72, slightly below logistic regression. Its sensitivity (true positive rate for defaulters) is 0.62, meaning it detects fewer defaulters compared to logistic regression. However, it achieves a marginally higher specificity, correctly clearing more non-defaulters.

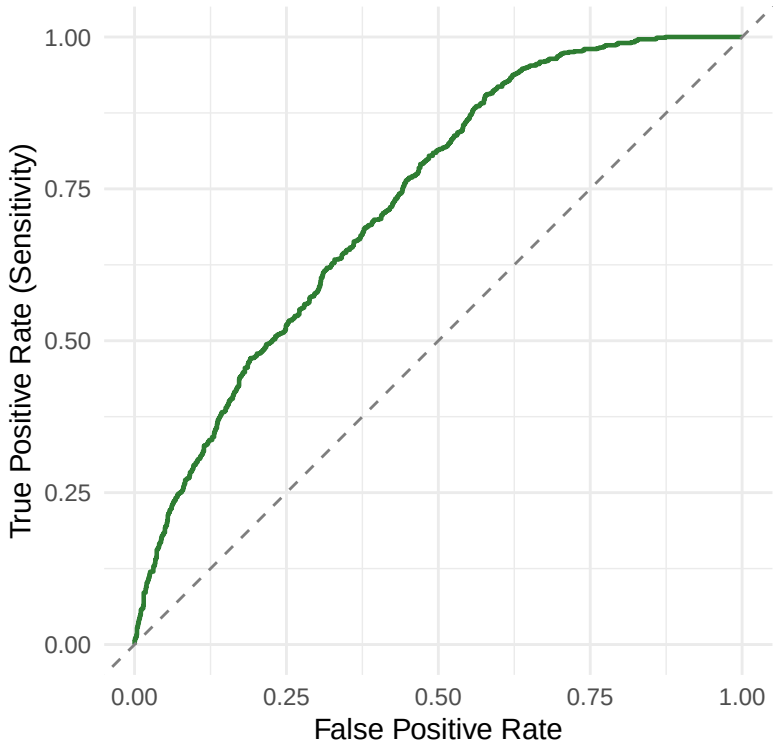

Figure 3: ROC Curve — k-NN (Test Set).

Table 5: Confusion Matrix — k-NN (Test Set)

	Predicted: Default	Predicted: No Default
Actual: Default	534	455
Actual: No Default	274	806

Interpretation: k-NN correctly identifies 502 defaulters (true positives) but misses 414 (false negatives). Compared to logistic regression, it catches fewer actual defaulters — a disadvantage in credit risk contexts where undetected defaults carry higher financial cost.

Section 4: Model Comparison

Table 6: Test Set Performance — Logistic Regression vs. k-NN

Model	ROC AUC	Accuracy	Sensitivity	Specificity
Logistic Regression	0.7439	0.6607	0.6473	0.6693
k-NN	0.7289	0.6477	0.6609	0.6392

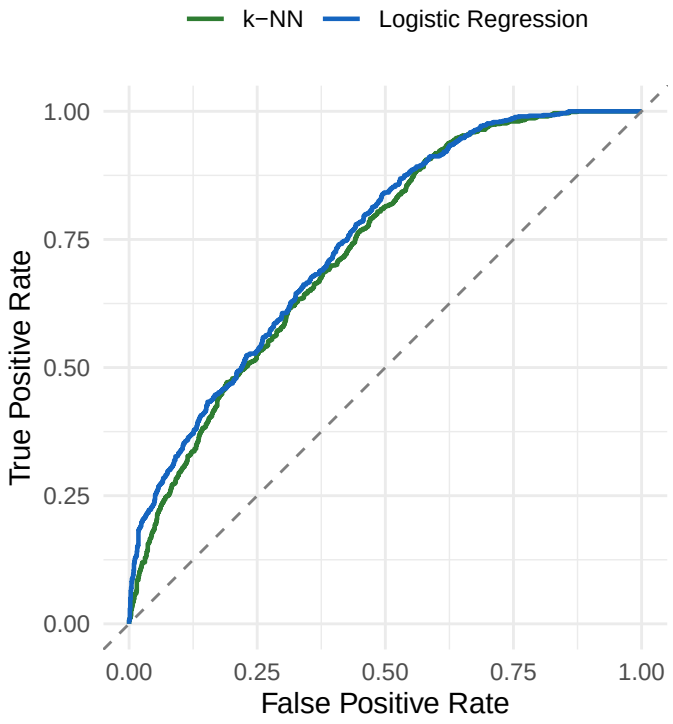

Figure 4: ROC Curves — Logistic Regression vs. k-NN (Test Set).

Interpretation: Logistic regression outperforms k-NN on ROC AUC (0.74 vs. 0.72), accuracy, and sensitivity. Both models achieve similar specificity. For BNPL default detection, logistic regression is the stronger performer of the two.

Section 5: Business Implications

- **Logistic regression** is the better-performing model of the two and has the added advantage of interpretability — each predictor's effect on default probability can be quantified and explained to risk managers and regulators.
- **k-NN** performs slightly below logistic regression on most metrics. Its lower sensitivity means more defaults go undetected, which is a significant drawback in a credit risk context.
- At the current 0.50 probability threshold, both models miss roughly a third of actual defaulters. A BNPL provider could lower the threshold to flag more potential defaulters, trading some specificity for higher sensitivity — a business decision depending on the relative cost of missed defaults versus false alarms.
- The predicted default probabilities from both models can support practical risk actions: prioritizing manual review for borderline cases, applying lower credit limits to high-risk customers, and triggering early repayment reminders.